

Problem Set 3: Calculus

Problem Set 3: Calculus

[Get the code](#)**Overview****Beta due** Thursday, April 6, 2017, 10:00 pm**Code reviews due** Monday, April 10, 2017, 11:00 am**Final due** Thursday, April 13, 2017, 10:00 pm**Problem 2: Parsing Expressions**

In this problem set, we will explore parsers, recursive data types, and equality for immutable types.

Compared to the previous problem sets, we are imposing very few restrictions on how you structure your code. In addition, much of the code that you write for problems in this problem set will depend heavily on how you decided to implement earlier parts of the problem set. We strongly recommend that you read through the entire assignment before writing any code.

Submitting**Design Freedom and Restrictions**

On several parts of this problem set, the classes and methods will be yours to specify and create, but you must pay attention to the **PS3 instructions** sections in the provided documentation.

You must satisfy the specifications of the provided interfaces and methods. You are, however, permitted to strengthen the provided specifications or add new methods. On this problem set, unlike previous problem sets, we will not be running your tests against any other implementations.

On this problem set, Didit provides less feedback about the correctness of your code:

- It is your responsibility to examine Didit feedback and make sure your code compiles and runs properly for grading.
- However, **correctness is your responsibility alone**, and you must rely on your own careful specification, testing, and implementation to achieve it.

Please remember to push early: we cannot guarantee timely Didit builds, especially near the problem set deadline.

Get the code

To get started,

1. Ask Didit to [create a remote psets/ps3 repository](#) for you on Athena.
2. Pull the repo from Athena using Git:

```
git clone ssh://[username]@athena.dialup.mit.edu/mit/6.031/git/sp17/psets/ps3/[username].git ps3
```

If you need a refresher on how to create, clone, and import your repository, see Problem Set 0.

Overview

Wouldn't it be nice not to have to differentiate all our calculus homework by hand every time? Wouldn't it be just lovely to type it into a computer and have the computer do it for us instead? For example, we could interact with it like this (user input in green):

If the output is an expression, your system may output an equivalent expression, including variations in spacing, parentheses, simplification, and number representation. Numbers must be precise.

```

> x * x * x
x*x*x

> !eval x=2
8

> !d/dx
(x*x)*1+(x*1+1*x)*x

> !eval x=0.5000
.75

> x * y
x*y

> !d/dy
0*y+x*1

```

In this system, a user can enter either an **expression** or a **command**.

An *expression* is a polynomial consisting of:

- addition +
- multiplication *
- exponentiation ^ to a literal nonnegative integer power (e.g. x^3)
- nonnegative numbers in decimal representation, which consist of digits and an optional decimal point (e.g. 7 and 4.2)
- variables, which are case-sensitive single-letters (e.g. y and F)
- parentheses (for grouping)

Order of operations uses the standard PEMDAS rules.

Space characters around symbols are irrelevant and ignored, so $2.3 * p$ means the same as $2.3 * p$. Spaces may not occur within numbers, so $2 . 3 * p$ is not a valid expression.

When the user enters an expression, that expression becomes the **current expression** and is echoed back to the user (user input in green):

```

> x * x * x
x*x*x

> x + x^2 * y + x
x + x^2*y + x

```

If the output is an expression, your system may output an equivalent expression, including variations in spacing, parentheses, simplification, and number representation. Numbers must be precise.

A *command* starts with !. The command operates on the current expression, and may also update the current expression. Valid commands:

d/dvar

produces the derivative of the current expression with respect to the variable *var*, and updates the current expression to that derivative

eval var₁=num₁ ... var_n=num_n

substitutes num_i for var_i in the current expression, and evaluates it to a single number if the expression contains no other variables; does **not** update the current expression

Entering an invalid expression or command prints an error but does not update the current expression. The error should include a human-readable message but is not otherwise specified.

More examples:

If the output is an expression, your system may output an equivalent expression, including variations in spacing, parentheses, simplification, and number representation. Numbers must be precise.

```

> x * x * x
x*x*x

> 3xy
ParseError: unknown expression

> !d/dx
(x*x)*1+(x*1+1*x)*x

> !d/dx
(x*x)*0+(x*1+1*x)*1+(x*1+1*x)*1+(x*0+1*1+x*0+1*1)*x

> !d/d
ParseError: missing variable in derivative command

> !eval
(x*x)*0+(x*1+1*x)*1+(x*1+1*x)*1+(x*0+1*1+x*0+1*1)*x

> !eval x=1
6.0

> !eval x=0 y=5
0

```

The three things that a user can do at the console correspond to three provided method specifications in the code for this problem set:

- `Expression.parse()`
- `Commands.differentiate()`
- `Commands.evaluate()`

These methods are used by `Main` to provide the console user interface described above.

Problem 1: we will create the `Expression` data type to represent expressions in the program.

Problem 2: we will create the parser that turns a string into an `Expression`, and implement `Expression.parse()`.

Problems 3-4: we will add new `Expression` operations for taking derivatives and evaluating, and implement `Commands.differentiate()` and `Commands.evaluate()`.

Problem 1: Representing Expressions

Define an immutable, recursive abstract data type to represent expressions as abstract syntax trees.

Your AST should be defined in the provided `Expression` interface (in `Expression.java`) and implemented by several concrete variants, one for each kind of expression. Each variant should be defined in its own appropriately-named `.java` file. You should have separate variant classes for `+`, `*`, and `^`.

Concrete syntax in the input, such as parentheses and whitespace, should not be represented at all in your AST.

1.1 Expression

To repeat, your data type must be **immutable** and **recursive**. Follow the recipe for creating an ADT:

- **Spec.** Choose and specify operations. For this part of the problem set, the only operations `Expression` needs are creators and producers for building up an expression, plus the standard observers `toString()`, `equals()`, and `hashCode()`. We are strengthening the specs for these standard methods; see below.
- **Test.** Partition and test your operations in `ExpressionTest.java`, including tests for `toString()`, `equals()`, and `hashCode()`. Note that we will not run your test cases on other implementations, just on yours.
- **Code.** Write the rep for your `Expression` as a data type definition in a comment inside `Expression`. Implement the variant classes of your data type.

Remember to include a Javadoc comment above every class and every method you write; define abstraction functions and rep invariants, and write `checkRep`; and document safety from rep exposure.

1.2 toString()

Define the `toString()` operation on `Expression` so it can output itself as a string. This string must be a valid expression as defined above. You have the freedom to decide how to format the output with whitespace and parentheses for readability, but the expression must have the same mathematical meaning.

Your `toString()` implementation must be recursive, and must not use `instanceof`.

Use the `@Override` annotation to ensure you are overriding the `toString()` inherited from `Object`.

Remember that your tests must obey the spec. If your `toString()` tests expect a certain formatting of whitespace and parentheses, you must specify this formatting in your spec.

1.3 equals() and hashCode()

Define the `equals()` and `hashCode()` operations on your AST to implement *structural equality*.

Structural equality defines two expressions to be equal if:

- the expressions contain the same variables, numbers, and operators;
- those variables, numbers, and operators are in the same order, read left-to-right;
- and they are grouped in the same way.

For example, the AST for $1 + x$ is *not* equal to the AST for $x + 1$, but it is equal to the ASTs for $1+x$, $(1+x)$, and $(1)+(x)$. The AST for x^2 is *not* equal to the AST for $x*x$.

For n -ary groupings where n is greater than 2:

- Such expressions must be equal to themselves. For example, the ASTs for $3 + 4 + 5$ and $(3 + 4 + 5)$ must be equal.
- However, whether they are equal or not to different groupings with the same mathematical meaning is *not specified*, and you should choose an appropriate specification and implementation for your AST. For example, you must determine whether the ASTs for $(3 + 4) + 5$ and $3 + (4 + 5)$ are equal.

Numbers must be equal if their decimal representations are equal. For example, the ASTs for $x + 1.5$ and $x + 1.5000$ must be equal. For numbers, you are required to handle nonnegative decimal numbers to arbitrary precision. Since the expression language only allows entering numbers in finite decimal representation, and then adding, multiplying, and raising them to a nonnegative integer power, every expression is guaranteed to produce a number that can be represented precisely in a finite decimal representation as well. It's a good idea to use the `BigDecimal` type for these numbers.

Remember: concrete syntax, including parentheses, should not be represented in your AST. Grouping, for example, should be reflected in the AST's structure.

Be sure that AST instances which are considered equal according to this definition and according to `equals()` also satisfy observational equality.

Your `equals()` and `hashCode()` implementations must be recursive. Only `equals()` can use `instanceof`, and `hashCode()` must not.

Remember to use the `@Override` annotation.

Commit to Git. Once you're happy with your solution to this problem, commit and push!

Problem 2: Parsing Expressions

Now we will create the parser that takes a string and produces an `Expression` value from it. The entry point for your parser should be `Expression.parse()`, whose spec is provided in the starting code.

Examples of valid inputs:

```

3 + 2.4
3 * x + 2.4
3 * (x + 2.4)
(3 * x) ^ 2
((3 + 4) * x * x)
a*x^2 + b*x + c
(2*x )+ ( y*x )
4 + 3 * x + 2 * x * x + 1 * x * x * ((x))

```

Examples of invalid inputs:

```

3 *  incomplete expression
( 3  incomplete expression
3 x  multiplication must use *
3^x  exponent must be nonnegative integer

```

Examples of optional inputs:

```

2 - 3
6.02e23
3^(2+1)
x/5

```

You may consider the optional inputs invalid, or you may choose to support additional features (new operators or number representations) in the input. However, *your system may not produce an output with a new feature unless that feature appeared in its input*. This way, a client who knows about your extensions can trigger them, but clients who don't know won't encounter them unexpectedly.

2.1 Write a grammar

Write a ParserLib grammar for polynomial expressions as described in the overview. A starting ParserLib grammar file can be found in `src/calculus/Expression.g`. This starting grammar recognizes sums of integers, and ignores spaces.

For more information on ParserLib, see:

- the reading on parser generators
- example code for compiling the parser and processing a parse tree
- Javadoc documentation for ParserLib

2.2 Implement `Expression.parse()`

Implement `Expression.parse()` by following the recipe:

- **Spec.** The spec for this method is given, but you may strengthen it if you want to make it easier to test.
- **Test.** Write tests for `Expression.parse()` and put them in `ExpressionTest.java`. Note that we will not run your tests on any implementations other than yours.
- **Code.** Implement `Expression.parse()` so that it calls the parser generated by your ParserLib grammar. The reading on parser generators discusses how to call the parser and construct an abstract syntax tree from it, including code examples.

2.3 Run the console interface

Now that `Expression` values can be both parsed from strings with `parse()`, and converted back to strings with `toString()`, you can try entering expressions into the console interface.

Run `Main`. In Eclipse, the Console view will allow you to type expressions and see the result. Try some of the expressions from the top of this handout.

Commit to Git. Once you're happy with your solution to this problem, commit and push!

Problem 3: Differentiation

The symbolic differentiation operation takes an expression and a variable, and produces an expression with the derivative of the input with respect to the variable. The result does not need to be simplified.

For example, the following are correct derivatives :

`x*x*x` with respect to `x`
`3 * x * x`

`x*x*x` with respect to `x`
`x*x + (x + x)*x`

`x*x*x` with respect to `x`
`((x*x) * 1) + ((x * 1) + (1 * x)) * x) + (0)`

`x^3` with respect to `x`
`3*x^2*1`

If the output is an expression, your system may output an equivalent expression, including variations in spacing, parentheses, simplification, and number representation. Numbers must be precise.

Incorrect derivatives:

`y*y*y` with respect to `y`
`0 d/dx, should be d/dy`

`x^0` with respect to `x`
`0*x^-1 no negative exponents`

Recall that expressions are polynomials, so the differentiation operation is well-defined, and its result can always be expressed as a polynomial as well.

To implement your recursive differentiation operation, use these rules:

$$\frac{dc}{dx} = 0 \quad \frac{dx}{dx} = 1 \quad \frac{d(u+v)}{dx} = \frac{du}{dx} + \frac{dv}{dx} \quad \frac{d(u \cdot v)}{dx} = u \left(\frac{dv}{dx} \right) + v \left(\frac{du}{dx} \right) \quad \frac{d}{dx} u^n = nu^{n-1} \frac{du}{dx}$$

where c is a constant or variable other than the variable we are differentiating with respect to (in this case x), n is a nonnegative integer literal, and u and v can be any expression, including x .

3.1. Add an operation to Expression

You should implement differentiation as a method on your `Expression` datatype, defined recursively. The signature and specification of the method are up to you to design. Follow the recipe:

- **Spec.** Define your operation in `Expression` and write a spec.
- **Test.** Put your tests in `ExpressionTest.java`. Note that we will not run your test cases on other implementations, just on yours.
- **Code.** The implementation must be recursive. It must not use `instanceof`, nor any equivalent operation you have defined that checks the type of a variant.

3.2 Implement Commands.differentiate()

In order to connect your differentiation operation to the user interface, we need to implement the `Commands.differentiate()` method.

- **Spec.** The spec for this operation is given, but you may strengthen it if you want to make it easier to test.
- **Test.** Write tests for `differentiate()` and put them in `CommandsTest.java`. These tests will likely be very similar to the tests you used for your lower-level differentiation operation, but they must use `Strings` instead of `Expression` objects. Note that we will not run your tests on any implementations other than yours.
- **Code.** Implement `differentiate()`. This should be straightforward: simply parsing the expression, calling your differentiation operation, and converting it back to a string.

3.3 Run the console interface

We've now implemented the `!d/d` command in the console interface. Run `Main` and try some derivatives in the Console view.

Commit to Git. Once you're happy with your solution to this problem, commit and push!

Problem 4: Evaluation

The evaluation operation takes an expression and an environment (a mapping of variables to values). It substitutes the values for those variables into the expression, and attempts to evaluate the substituted polynomial as much as it can.

The set of variables in the environment is allowed to be different than the set of variables actually found in the expression. Any variables in the expression but not the environment remain as variables in the substituted polynomial. Any variables in the environment but not the expression are simply ignored.

The only required evaluation is that if the substituted polynomial is a constant expression, with no variables remaining, then evaluation must reduce it to a single number, with no operators remaining either. Evaluation for substituted polynomials that still contain variables is underdetermined, left to the implementer's discretion. You may strengthen this spec if you wish to require particular evaluations in other cases.

For example, the following are correct output for simplified expressions:

`x*x*x` with environment `x=5, y=10, z=20`
125

`x*x*x + y^3` with environment `y=10`
`x*x*x+1000`

`x*x*x + y^3` with environment `y=10`
`x*x*x+10*10*10`

`1+2^3+8*0.5` with empty environment
13.000

If the output is an expression, your system may output an equivalent expression, including variations in spacing, parentheses, simplification, and number representation. Numbers must be precise.

Incorrect simplified expressions:

`x*x*y` with environment `x=1, y=3`
`1*1*3` *not a single number*

`x^3` with environment `x=2`
(8) *includes parentheses*

Optional simplified expressions:

`x*x*y + y*(1+x)` with environment `x=2`
`7*y`

`x*x*x + x*x*x` with empty environment
`2*x^3`

4.1 Add an operation to Expression

You should implement evaluation as a method on your `Expression` datatype, defined recursively. The signature and specification of the method are up to you to design. Follow the recipe:

- **Spec.** Define your operation in `Expression` and write a spec.
- **Test.** Put your tests in `ExpressionTest.java`. Note that we will not run your test cases on other implementations, just on yours.
- **Code.** The implementation must be recursive (perhaps by calling recursive helper methods). It must not use `instanceof`, nor any equivalent operation you have defined that checks the type of a variant class.

You may find it useful to add more operations to `Expression` to help you implement the evaluate operation. Spec/test/code them using the same recipe, and make them recursive as well where appropriate. Your helper operations should not simply be a variation on using `instanceof` to test for a variant class.

You may also find the `Optional` class useful, as a typesafe alternative for a situation where you might be tempted to return `null`.

4.2 Commands.evaluate()

In order to connect your evaluate operation to the user interface, we need to implement the `Commands.evaluate()` method.

- **Spec.** The spec for this operation is given, but you may strengthen it if you want to make it easier to test.

- **Test.** Write tests for `evaluate()` and put them in `CommandsTest.java`. These tests will likely be very similar to the tests you used for your lower-level `evaluate` operation, but they should use `Strings` instead of `Expression` objects. Note that we will not run your tests on any implementations other than yours.
- **Code.** Implement `evaluate()`. This should be straightforward: simply parsing the expression, calling your `evaluate` operation, and converting it back to a string.

4.3 Run the console interface

We've now implemented the `!eval` command in the console interface. Run `Main` and try using it in the Console view.

Commit to Git. Once you're happy with your solution to this problem, commit and push!

Before you're done

- Make sure you have **documented specifications**, in the form of properly-formatted Javadoc comments, for all your types and operations.
- Make sure you have **documented abstraction functions and representation invariants**, in the form of a comment near the field declarations, for all your implementations.

With the rep invariant, also say how the type prevents rep exposure.

Make sure all types use `checkRep()` to check the rep invariant and implement `toString()` with a useful representation of the abstract value.

- Make sure you have satisfied the **Object contract** for all types. In particular, you will need to specify, test, and implement `equals()` and `hashCode()` for all immutable types.
- Use `@Override` when you override `toString()`, `equals()`, and `hashCode()`, to gain static checking of the correct signature.

Also use `@Override` when a class implements an interface method, to remind readers where they can find the spec.

- Make sure you have a thorough, principled **test suite** for every type. Note that `Expression`'s variant classes are considered part of its rep, so a single good test suite for `Expression` covers the variants too.

Submitting

Make sure you commit AND push your work to your repository on Athena. We will use the state of your repository on Athena as of 10:00pm on the deadline date. When you `git push`, the continuous build system attempts to compile your code and run some basic tests. You can always review your build results at didit.csail.mit.edu.

Didit feedback is provided on a best-effort basis:

- There is no guarantee that Didit tests will run within any particular timeframe, or at all. If you push code close to the deadline, the large number of submissions will slow the turnaround time before your code is examined.
- If you commit and push right before the deadline, the Didit build does not have to complete in order for that commit to be graded.
- Passing some or all of the **public tests** on Didit is no guarantee that you will pass the full battery of autograding tests — but failing them is almost sure to mean lost points on the problem set.

MIT EECS

